

# ICPS Anomaly Detection

## Project Configuration

*Samruddhi Khairnar*

### Introduction :

This project is aimed at building a simulated industrial cyber physical system using the Electra dataset, training an SVM model on it for predicting anomalous points in the dataset, processing and predicting the labels of a simulated stream of incoming data points and showing the output in the form of a dashboard.

### Technology Stack :

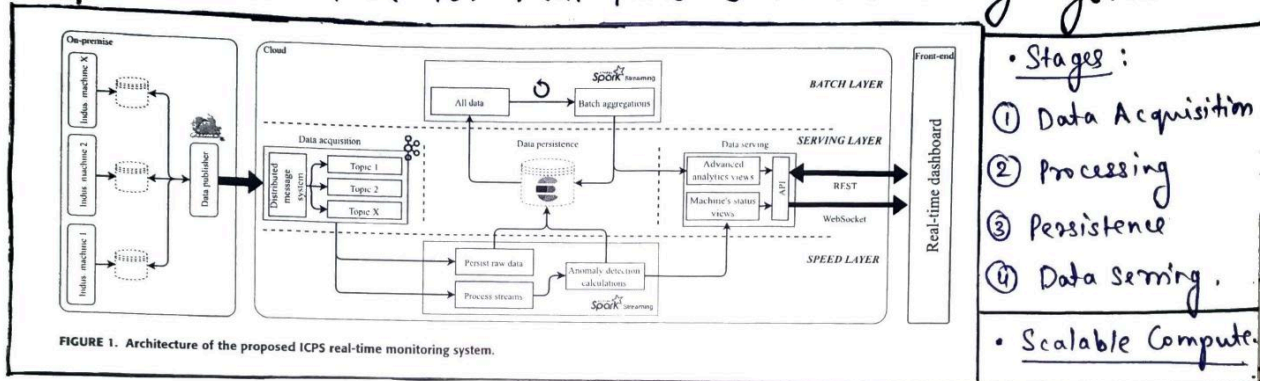
- Python3, JupyterLab, Colab.
- Dash library for Dashboard.
- Apache Spark for distributed ML model (SVM) training.
- Apache Hadoop for distributed storage.
- Apache Kafka for simulated pub-sub producer-consumer model for streaming data points in real time.

### References :

1. **Project Architecture :**  
<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8693711>
2. **Electra Dataset :**  
<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8926471>

## Proposed Architecture (as mentioned in the base paper) :

### Proposed Architecture for Real-Time ICPS Monitoring System.



#### • Stages :

- ① Data Acquisition
- ② Processing
- ③ Persistence
- ④ Data serving.

#### • Scalable Compute.

- Layers : ① Kafka for Data Acquisition : publish - subscribe model + Flume.
- ② Realtime + Batch processing using Spark Streaming (detect Anomalies).
- ③ Persistence : Distributed storage + querying : Elastic Search + Influx DB.
- ④ Zookeeper : Cloud Resource Management.
- ⑤ Serving : Rest API : Querying + Websocket : Dashboard + communication.

- Detection : ① Single Data Point (SDAD)
- ② Multiple /Batch (MDAD).

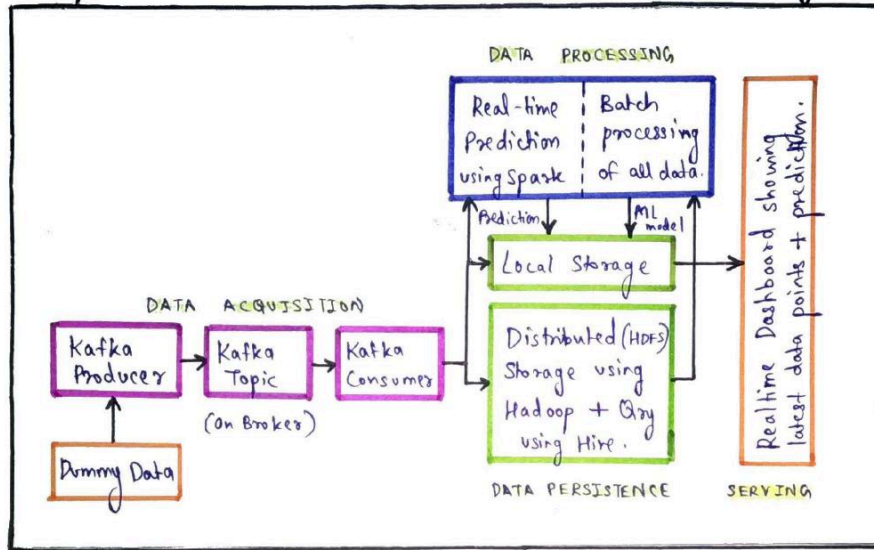
- Evaluation : • Input rate • Scheduling delay
- (Scale tests) • Processing time • Delay Time.

- Future : Predict early repair work to increase OEE (Overall Equipment Effectiveness)

- Reference : 10.1109/ACCESS.2019.2911979

## Implemented Architecture (using a 3-node cluster on AWS EC2) :

### Implemented Architecture for ICPS Anomaly Detection System.



#### • Layers:

##### ① Data acquisition:

A stream of data points is simulated using Kafka's producer and consumer.

##### ② Persistence:

The data points are stored in HDFS, in a

distributed manner. ③ Data Processing: The stored data is batch-processed using Apache Spark. Also, the output of it: A trained ML model is saved for predicting the labels of the incoming stream of data points.

④ Data Serving: The latest data points are visualized using a Dashboard.

- Tech Stack:
- Hadoop (HDFS, YARN, Hive, Zookeeper)
  - Spark
  - PySpark
  - Kafka
  - Python 3
  - Plotly-Dash
  - Colab / Jupyter NB.
  - Bash
  - Hosted on AWS EC2 virtual machines.

## Electra Dataset Description (as mentioned in the reference paper) :

Dataset Used : Electra-Modbus Dataset (12 hr traffic captured) - 4 nodes.

- Features:
  - ① time : timestamp of packet.
  - ② smac : Source MAC address.
  - ③ dmac : Destination MAC address.
  - ④ sip : Source IP address.
  - ⑤ dip : Destination IP address.
  - ⑥ request : 1 = master to slave comm<sup>n</sup>.
  - ⑦ fc : function code. (Read holding register / Read input ... others = anomalies).
  - ⑧ address : Memory address to perform Read/Write operation.
  - ⑨ error : Indicates an error in Read/Write operation.
  - ⑩ Label : Labels for attacks/normal samples.
  - ⑪ data : Read/Written data.

- Labels:
  - ① Normal
  - ② MITM-UNALTERED  $\approx$  Normal (Man in the Middle node is active but does no change).
  - ③ READ-ATTACK :  $fc = 2$  - packet generation.
  - ④ WRITE-ATTACK : Wrong data packet generated.
  - ⑤ FUNCTION-CODE-RECOGNITION-ATTACK : Generating wrong  $fc$  packets.
  - ⑥ REPLAY-ATTACK : Replaying normal packets at wrong time intervals.
  - ⑦ RESPONSE-MODIFICATION-ATTACK : Changing data inside a normal packet. (IP changes).
  - ⑧ FORCE-ERROR-ATTACK : Modifying error field of normal packet.

- Experiments:
  - ① Svm : 99% precision score.
  - ② Random Forest : 97%.
  - ③ Neural Network : 94%.

• Reference: 10.1109, ACCESS, 2019. 2958284.

## Future Scope :

Future Scope :

- ① Real-time push notifications / alerts.
- ② Auto-scaling of cluster size.
- ③ Using advanced deep learning models for batch-processing.
- ④ Automating the entire pipeline using scripts.
- ⑤ Deploying the system in actual industrial system use cases.




## Dashboard :

```
tweety@tweety-HP-ProBook-450-G3: ~  
File Edit View Search Terminal Help  
tweety@tweety-HP-ProBook-450-G3:~$ start-all.sh  
WARNING: Attempting to start all Apache Hadoop daemons as tweety in 10 seconds.  
WARNING: This is not a recommended production deployment configuration.  
WARNING: Use CTRL-C to abort.  
Starting namenodes on [localhost]  
Starting datanodes  
Starting secondary namenodes [tweety-HP-ProBook-450-G3]  
Starting resourcemanager  
Starting nodemanagers  
tweety@tweety-HP-ProBook-450-G3:~$ jps  
15348 ResourceManager  
15494 NodeManager  
15015 Worker  
14855 SecondaryNameNode  
15784 QuorumPeerMain  
tweety@tweety-HP-ProBook-450-G3:~$  
tweety@tweety-HP-ProBook-450-G3:~/kafka_2.13-3.4.0$ bin/zookeeper-server-start.sh config/zookeeper.properties  
[2023-07-07 09:24:24,470] INFO Reading configuration from: config/zookeeper.properties (org.apache.zookeeper.server.quorum.QuorumPeerConfig)  
[2023-07-07 09:24:24,490] WARN config/zookeeper.properties is relative. Prepend ./ to indicate that you're sure! (org.apache.zookeeper.server.quorum.QuorumPeerConfig)  
[2023-07-07 09:24:24,524] INFO clientPortAddress is 0.0.0.0:2181 (org.apache.zookeeper.server.quorum.QuorumPeerConfig)  
[2023-07-07 09:24:24,525] INFO serverPortAddress is not set (org.apache.zookeeper.server.quorum.QuorumPeerConfig)  
tweety@tweety-HP-ProBook-450-G3:~/kafka_2.13-3.4.0$ bin/kafka-server-start.sh config/server.properties  
[2023-07-07 09:25:31,282] INFO Registered kafka:type=kafka.Log4jController MBean (kafka.tools.Log4jControllerRegistration$)  
[2023-07-07 09:25:33,091] INFO Setting -D jdk.tls.rejectClientInitiatedRenegotiation=true (org.apache.kafka.server.KafkaServer$)  
tweety@tweety-HP-ProBook-450-G3:~$ python3 dashboard.py  
Dash is running on http://0.0.0.0:8050/  
* Serving Flask app 'dashboard'  
* Debug mode: off  
tweety@tweety-HP-ProBook-450-G3:~$ python3 producer.py  
tweety@tweety-HP-ProBook-450-G3:~$ python3 consumer.py  
[2023-07-07 09:29:40] WARN Utils: Your hostname, tweety-HP-ProBook-450-G3, is not set (org.apache.kafka.common.utils.Utils)
```

09:43:51  
Samruddhi K

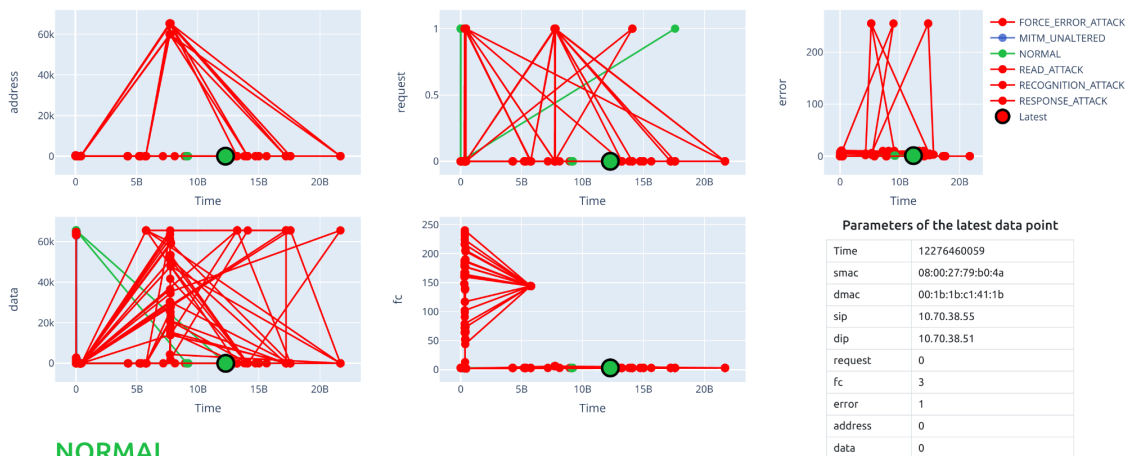
### Anomaly Detection in ICPS simulated using the Electra Dataset



READ\_ATTACK

09:51:18  
Samruddhi K

### Anomaly Detection in ICPS simulated using the Electra Dataset



NORMAL

## Original Source Code for Distributed 3-node cluster :

1. Configuration : [bit.ly/icps-1](http://bit.ly/icps-1)
2. Exploratory Data Analysis : [bit.ly/icps-2](http://bit.ly/icps-2)
3. Model Training : [bit.ly/icps-3](http://bit.ly/icps-3)

## Setup for Standalone single-node cluster :

### Prerequisites

1. Install Ubuntu.
2. Install curl :

```
tweety@tweety-HP-ProBook-450-G3:~$ sudo apt-get update
tweety@tweety-HP-ProBook-450-G3:~$ sudo apt install curl
```

3. Install Java 11:

```
tweety@tweety-HP-ProBook-450-G3:~$ sudo apt install openjdk-11-jdk
```

4. Check whether Python3 is installed :

```
tweety@tweety-HP-ProBook-450-G3:~$ python3 --version
```

5. Install gdown + add PATH :

```
tweety@tweety-HP-ProBook-450-G3:~$ pip install gdown
tweety@tweety-HP-ProBook-450-G3:~$ export PATH=/home/<username>/.local/bin
# Replace with your username
```

### Hadoop, Spark, Kafka Installation

1. Download the configured setups + untar it (setup for single node cluster) :

```
tweety@tweety-HP-ProBook-450-G3:~$ gdown 10byFo7JAgABKUUD0kcGBZ-RgoPQ51q_D
tweety@tweety-HP-ProBook-450-G3:~$ gdown 1UJQDaSWbAXThG4jp9eVjbESI7gISd0Dh
tweety@tweety-HP-ProBook-450-G3:~$ gdown 1QoMG2H4EBVtM09rQtY4GZFz65FT-XeKh
tweety@tweety-HP-ProBook-450-G3:~$ gdown 1hbUdWQso2xYTt70A4-PWFIDiJOfxS9q
tweety@tweety-HP-ProBook-450-G3:~$ gdown 1P7k3LLt1QHrIdewdtYghfFeQdviwg-gw
tweety@tweety-HP-ProBook-450-G3:~$ gdown 1S5snu1xmB4Arcyf36ICu4n-HqID-0a7a
tweety@tweety-HP-ProBook-450-G3:~$ tar -xzvf hadoop.tar.gz
tweety@tweety-HP-ProBook-450-G3:~$ tar -xzvf spark.tar.gz
tweety@tweety-HP-ProBook-450-G3:~$ tar -xzvf kafka.tar.gz
```

2. Make these temp directories (as configured in the setup ./hadoop-3.3.5/etc/hadoop/core-site and hdfs-site.xml + required for kafka, zookeeper logs) + give them sudo privileges :

```
tweety@tweety-HP-ProBook-450-G3:~$ sudo mkdir /hadoop
tweety@tweety-HP-ProBook-450-G3:~$ sudo chmod -R 777 /hadoop/
tweety@tweety-HP-ProBook-450-G3:~$ sudo chown -R <username> /hadoop
# Replace with your username

tweety@tweety-HP-ProBook-450-G3:~$ sudo mkdir /hadoop/tmp
tweety@tweety-HP-ProBook-450-G3:~$ sudo mkdir /hadoop/dfs
tweety@tweety-HP-ProBook-450-G3:~$ sudo mkdir /hadoop/dfs/data
tweety@tweety-HP-ProBook-450-G3:~$ sudo mkdir /hadoop/dfs/name
tweety@tweety-HP-ProBook-450-G3:~$ sudo mkdir /app
tweety@tweety-HP-ProBook-450-G3:~$ sudo mkdir /app/zookeeper
tweety@tweety-HP-ProBook-450-G3:~$ sudo chmod 777 /app/zookeeper
tweety@tweety-HP-ProBook-450-G3:~$ sudo mkdir /app/kafka-logs
tweety@tweety-HP-ProBook-450-G3:~$ sudo mkdir /app/kafka-logs/logs
tweety@tweety-HP-ProBook-450-G3:~$ sudo chmod 777 /app/kafka-logs/logs
```

3. Get the present working directory and copy the path (ctrl+shift+C) :

```
tweety@tweety-HP-ProBook-450-G3:~$ pwd
```

4. Add the installation directories to the PATH, open bashrc :

```
tweety@tweety-HP-ProBook-450-G3:~$ nano .bashrc
```

Type these lines to end of the file then, add pwd copied from previous step and (save : ctrl+S -> exit : ctrl+X) :

```
export HADOOP_HOME=<pwd>/hadoop-3.3.5/
export JAVA_HOME="/usr/lib/jvm/java-11-openjdk-amd64/"
export SPARK_HOME=<pwd>/spark-3.4.0-bin-hadoop3/
export KAFKA_HOME=<pwd>/kafka_2.13-3.4.0/
export PATH=$PATH:$HADOOP_HOME/bin
export PATH=$PATH:$HADOOP_HOME/sbin
```

5. Source the changes :

```
tweety@tweety-HP-ProBook-450-G3:~$ source .bashrc
```

## Source Code

1. Unzip and move the source code files :

```
tweety@tweety-HP-ProBook-450-G3:~$ sudo apt install unzip
tweety@tweety-HP-ProBook-450-G3:~$ export PATH=/usr/bin:/bin
```

```
tweety@tweety-HP-ProBook-450-G3:~$ unzip electra-model.zip
tweety@tweety-HP-ProBook-450-G3:~$ unzip files.zip
tweety@tweety-HP-ProBook-450-G3:~$ unzip files2.zip
tweety@tweety-HP-ProBook-450-G3:~$ mv files/* . -f
tweety@tweety-HP-ProBook-450-G3:~$ mv files2/* . -f
```

## Installation of Dependencies

1. Install the dependencies :

```
tweety@tweety-HP-ProBook-450-G3:~$ pip3 install pandas pyspark confluent_kafka dash dash_bootstrap_components
```

## Start Hadoop, Spark, Kafka

1. Setup passwordless SSH :

```
tweety@tweety-HP-ProBook-450-G3:~$ sudo apt install openssh-client
openssh-server
```

```
# Add the following lines to /etc/ssh/sshd_config
```

```
PubkeyAuthentication yes
```

```
PasswordAuthentication no
```

```
ChallengeResponseAuthentication no
```

```
tweety@tweety-HP-ProBook-450-G3:~$ chmod 700 .ssh | chmod 644
.ssh/id_rsa.pub | chmod 600 .ssh/authorized_keys | chmod 600 .ssh/id_rsa
| chmod 755 ~
```

```
# Init ssh each time using :
```

```
sudo service ssh restart (or start)
```

```
tweety@tweety-HP-ProBook-450-G3:~$ ssh-keygen -t rsa; cp
~/.ssh/id_rsa.pub ~/.ssh/authorized_keys
```

```
# Press enter for each line
```

```
# Verify the passwordless setup using :
```

```
ssh localhost
```

2. Start Hadoop and put the initial dataset in the HDFS :

```
tweety@tweety-HP-ProBook-450-G3:~$ source .bashrc
tweety@tweety-HP-ProBook-450-G3:~$ hdfs namenode -format
tweety@tweety-HP-ProBook-450-G3:~$ start-all.sh
```



```
# Verify the running of Hadoop using :
tweety@tweety-HP-ProBook-450-G3:~$ jps

tweety@tweety-HP-ProBook-450-G3:~$ hdfs dfs -mkdir /electra
tweety@tweety-HP-ProBook-450-G3:~$ hdfs dfs -put ./d.csv /electra/electra_modbus.csv
```

### 3. Start Spark :

```
tweety@tweety-HP-ProBook-450-G3:~$ $SPARK_HOME/sbin/start-all.sh
```

### 4. Install and start Tilix - *will be required for multiple terminal windows simultaneously (example shown below)* - use the buttons to open multiple terminals :

```
tweety@tweety-HP-ProBook-450-G3:~$ sudo apt install tilix
tweety@tweety-HP-ProBook-450-G3:~$ tilix
```



### 5. In one Tilix terminal, start Zookeeper (*keep it open*):

```
tweety@tweety-HP-ProBook-450-G3:~$ cd $KAFKA_HOME
tweety@tweety-HP-ProBook-450-G3:~$ bin/zookeeper-server-start.sh
config/zookeeper.properties
```

### 6. In another Tilix terminal and start Kafka server (*keep it open*) :

```
tweety@tweety-HP-ProBook-450-G3:~$ cd $KAFKA_HOME
tweety@tweety-HP-ProBook-450-G3:~$ bin/kafka-server-start.sh
config/server.properties
```

### 7. In the third Tilix terminal, create a new Kafka topic :

```
tweety@tweety-HP-ProBook-450-G3:~$ cd $KAFKA_HOME/bin
tweety@tweety-HP-ProBook-450-G3:~$ ./kafka-topics.sh --create --topic
electra --bootstrap-server localhost:9092
```

8. In the fourth Tilix terminal, start the Kafka consumer (*consumes from topic electra and stores in HDFS + locally in d.csv*) :

```
tweety@tweety-HP-ProBook-450-G3:~$ cd
tweety@tweety-HP-ProBook-450-G3:~$ mv ./content/electra-model/ .
tweety@tweety-HP-ProBook-450-G3:~$ python3 consumer.py
```

9. In the fifth Tilix terminal, start the Kafka producer (*pushes data into topic electra and from dummy.csv*) :

```
tweety@tweety-HP-ProBook-450-G3:~$ python3 producer.py
```

10. In the sixth Tilix terminal, start the dashboard process (*displays from d.csv - latest 200 data points and gives params + prediction of the latest data point*) :

```
tweety@tweety-HP-ProBook-450-G3:~$ python3 dashboard.py
```

11. Run ml.py once, every week, to train on the entire dataset :

*Create new file to store the cron job command, make it executable and open it :*

```
tweety@tweety-HP-ProBook-450-G3:~$ touch run.sh
tweety@tweety-HP-ProBook-450-G3:~$ chmod +x cron.sh
tweety@tweety-HP-ProBook-450-G3:~$ nano run.sh
```

*Enter this in to the file (paste : ctrl+shift+V -> save : ctrl+S -> exit : ctrl+X) :*

```
#!/bin/bash
while true; do
python3 ml.py
sleep 604800
done
```

Run this cronjob in the seventh Tilix Terminal :

```
tweety@tweety-HP-ProBook-450-G3:~$ ./run.sh
```

**View the Dashboard using a Browser :**

<http://127.0.0.1:8050/>

## Stop all Processes :

Stop Hadoop :

```
tweety@tweety-HP-ProBook-450-G3:~$ stop-all.sh
```

Stop Spark :

```
tweety@tweety-HP-ProBook-450-G3:~$ $SPARK_HOME/sbin/stop-all.sh
```

Stop other processes running in terminals (Kafka, Zookeeper, Python3 codes) :

```
Press Ctrl + C
```